

Differentiate Colorectal carcinoma and para-carcinoma tissue based on expression profile

Title: Hypoxia-induced cysteine metabolism reprogramming are crucial for the tumorigenesis of colorectal cancer

Gene Expression Omnibus: GSE223119

Data Analyzer: Kai-Wei Chang

Data Owner: Shiyi Yang

Email: octoberfirst@sjtu.edu.cn

Affiliation: Shanghai Jiao Tong University, Shanghai, China

Q1: Main objective of the analysis that specifies whether your model will be focused on prediction or interpretation and the benefits that your analysis provides to the business or stakeholders of this data.

Answer:

This analysis will be focused on prediction

By determining whether the gene expression profiles capable of classifying the sample's origin, future expression profiles of unknown origin may be classified accordingly. In addition, the analysis will compare and contrast the performance of classification tools on low sample, high-dimension data set.

ie. each gene is assumed to be independent feature for this analysis

Q2: Brief description of the data set you chose, a summary of its attributes, and an outline of what you are trying to accomplish with this analysis.

Answer:

Metabolic reprogramming is a hallmark of human cancer and cancer-specific metabolism provide opportunities for cancer diagnosis, prognosis, and treatment. However, how metabolic pathways affect the initiation and progression of colorectal cancer remain largely unknown.

This data set includes 40 gene expression profiles of cancer and para-carcinoma tissue (tissue surrounding cancer).

This analysis attempt to build a model for predicting the tissue origin of unknown gene profiles

```
In [1]: ##No warning output
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn
```

Q3: Brief summary of data exploration and actions taken for data cleaning and feature engineering.

Answer:

Data Exploration:

- Raw data is composed of gene ID, gene expression on each sample, gene names, and database and gene network category
- Gene expression has been preprocessed by the data provider, with no missing values
- This data set contain extremely high amount of features (ie. 61,700 genes)

Data Engineering:

- For the project purpose, only gene ID (feature), gene expression, and tissue origin will be considered.
- Removing features that are not expressing or have very low variability among all profiles
- For simplification in this analysis, I randomly selected 3000 features for analysis
- Gene expression may or may not be skewed, and expression levels vary from gene to gene, standardization is applied

Detail analysis see below

```
In [2]: #Import Data processing libraries
import numpy as np
import pandas as pd
```

```
In [3]: #Load Data, data is preprocessed
rawData=pd.read_csv(r"C:\Users\kai-w\Desktop\03_Supervised Machine Learning Classification\GSE223119_MJ20190424016-gene.tpm.matr.ix.annot.txt', sep='\t', header=0)
```

```
In [4]: #Data overview
rawData.head()
```

	Unnamed: 0	C10	C11	C12	C13	C14	C15	C18	C1	C20	...	cog	cog_description	KO_id	KO_name	paths	pfam	go	n
0	ENSG000000000003	61.65	37.27	85.58	68.76	32.17	49.63	64.18	10.96	36.86	...	ENOG4111IRYS(Function unknown)	ENOG4111IRY(Tenomodulin)	K17295	TSPAN6	NaN	PF00335.17(Tetraspannin:Tetraspanin family)	GO:0039532(biological_process:negative regulat...	NP_003261.1(tetraspanin-I isoform a [Homo sapi.
1	ENSG000000000005	0.31	3.39	1.34	2.24	0.25	4.48	1.64	0.15	0.89	...	ENOG410YB96S(Function unknown)	ENOG410YB96(Tenomodulin)	NaN	NaN	NaN	PF04089.11(BRICHOS:BRICHOS domain)	GO:0005737(cellular_component:cytoplasm); GO:0...	XP_006986474.1(PREDICTED tenomodulin [Peromys.
2	ENSG0000000000419	81.96	34.78	56.60	81.13	49.78	121.31	91.79	39.23	41.02	...	COG0463(M:Cell wall/membrane/envelope biogenesis)	COG0463(Glycosyl transferase, family 2)	K00721	DPM1	map00510(N-Glycan biosynthesis)	PF00535.23(Glycos_transf_2:Glycosyl transferas...	GO:0019673(biological_process:GDP-mannose meta...	NP_001303964.1(dolichol phosphate mannosyltran.
3	ENSG0000000000457	4.98	1.42	2.84	4.01	4.90	3.86	3.63	4.08	3.95	...	ENOG410XQTGS(Function unknown)	ENOG410XQTGS(cerevisiae)	K17542	SCYL3	NaN	PF00069.22(Pkinase:Protein kinase domain); PF0...	GO:0005794(cellular_component:Golgi apparatus)...	XP_003893590.2(protein associating with the ca.
4	ENSG0000000000460	7.14	1.99	2.80	9.40	5.90	4.13	9.92	4.05	4.21	...	ENOG4110VTC(S(Function unknown)	ENOG4110VTC(Chromosome 1 open reading frame 112)	NaN	NaN	NaN	PF14868.3(DUF4487:Domain of unknown function [...	NaN	XP_005245374.1(uncharacterize protein C1orf11.

5 rows x 54 columns

```
In [5]: #Data types for each column
rawData.dtypes
```

```
Out[5]: Unnamed: 0      object
C10      float64
C11      float64
C12      float64
C13      float64
C14      float64
C15      float64
C18      float64
C1       float64
C20      float64
C22      float64
C23      float64
C24      float64
C25      float64
C3       float64
C4       float64
C5       float64
C6       float64
C7       float64
C8       float64
C9       float64
P10      float64
P11      float64
P12      float64
P13      float64
P14      float64
P15      float64
P18      float64
P1       float64
P20      float64
P22      float64
P23      float64
P24      float64
P25      float64
P3       float64
P4       float64
P5       float64
P6       float64
P7       float64
P8       float64
P9       float64
gene_name object
length    int64
description object
cog        object
cog_description object
K0_id      object
K0_name    object
paths      object
pfam       object
go         object
nr         object
swissprot  object
entrez     float64
dtype: object

In [6]: #Transpose so treat genes as feature, and each profile as independent record
data=rawData.loc[:,rawData.columns.str.match('^C|^P')].rename(index=rawData['Unnamed: 0']).T
data.head()
```

	ENSG000000000003	ENSG000000000005	ENSG000000000419	ENSG000000000457	ENSG000000000460	ENSG000000000938	ENSG000000000971	ENSG00000001036	ENSG00000001084	ENSG00000001167	...	MSTRG.9662	MSTRG.9664	MSTRG.9666	MSTRG.9667	MSTRG.9925
C10	61.65	0.31	81.96	4.98	7.14	4.40	15.56	38.45	18.87	12.65	...	0.02	0.00	0.0	0.13	0.37
C11	37.27	3.39	34.78	1.42	1.99	3.10	5.15	34.72	11.98	3.12	...	0.10	6.57	0.0	1.17	0.61
C12	85.58	1.34	56.60	2.84	2.80	5.31	8.48	29.41	25.83	11.12	...	0.07	0.00	0.0	0.20	0.17
C13	68.76	2.24	81.13	4.01	9.40	4.46	14.99	45.74	25.93	13.76	...	0.04	0.00	0.0	0.33	0.75
C14	32.17	0.25	49.78	4.90	5.90	2.46	10.40	33.28	17.18	11.91	...	0.27	7.07	0.0	0.36	0.54

5 rows × 61700 columns

```
In [7]: #Assign tissue types for classification
data['Tissue']=data.index.str[0:1]
data.head()
```

	ENSG000000000003	ENSG000000000005	ENSG000000000419	ENSG000000000457	ENSG000000000460	ENSG000000000938	ENSG000000000971	ENSG00000001036	ENSG00000001084	ENSG00000001167	...	MSTRG.9664	MSTRG.9666	MSTRG.9667	MSTRG.9925	MSTRG.9927
C10	61.65	0.31	81.96	4.98	7.14	4.40	15.56	38.45	18.87	12.65	...	0.00	0.0	0.13	0.37	0.00
C11	37.27	3.39	34.78	1.42	1.99	3.10	5.15	34.72	11.98	3.12	...	6.57	0.0	1.17	0.61	1.78
C12	85.58	1.34	56.60	2.84	2.80	5.31	8.48	29.41	25.83	11.12	...	0.00	0.0	0.20	0.17	0.50
C13	68.76	2.24	81.13	4.01	9.40	4.46	14.99	45.74	25.93	13.76	...	0.00	0.0	0.33	0.75	0.23
C14	32.17	0.25	49.78	4.90	5.90	2.46	10.40	33.28	17.18	11.91	...	7.07	0.0	0.36	0.54	0.20

5 rows × 61701 columns

```
In [8]: #check for presence of empty values
data.isna().values.sum()

Out[8]: 0

In [9]: #huge feature size, reduce by taking randomly 5000 for the purpose of the project
data.shape

Out[9]: (40, 61701)

In [10]: #some features/genes have low variability, remove those
feature_var=pd.DataFrame([([1, data[i].std()] for i in data.columns[:-1]),
                           columns=['feature', 'std']].set_index('feature'))
print(feature_var[feature_var['std']>0].sort_values('std'))

feature
ENSG000000232264    0.001581
ENSG000000228856    0.001581
ENSG000000231051    0.001581
ENSG000000133105    0.001581
ENSG000000156925    0.001581
...
ENSG000000211890    6428.959034
ENSG000000240040    6780.842254
ENSG000000212907    7435.193295
ENSG000000198804    8216.179258
ENSG000000228253    18410.117553

[41928 rows x 1 columns]

In [11]: #Random select 3000 features/genes that are variable for the project's purpose
features=feature_var[feature_var['std']>0.5].index.to_series().sample(3000).to_list()
data[features].head()
```


	ENSG000000283050	ENSG000000178971	ENSG0000000267530	ENSG000000181274	ENSG000000137414	ENSG000000268621	ENSG000000175309	ENSG000000103047	ENSG000000129226	ENSG000000133818	...	ENSG000000181754	MSTRG.35998	ENSG000000054654	ENSG00000001120
C10	3.51	7.54	0.81	33.93	12.00	0.08	38.41	1.90	61.67	11.94	...	2.31	1.91	22.00	5
C11	0.82	8.99	1.85	18.74	4.18	129.29	12.25	2.26	111.31	9.39	...	0.65	12.09	13.21	3
C12	2.89	8.23	0.99	36.07	5.21	0.42	19.09	2.71	62.22	12.51	...	2.01	5.18	36.34	2
C13	3.47	5.18	0.41	18.03	8.45	0.00	30.69	4.56	64.74	19.50	...	1.16	4.44	20.36	5
C14	2.82	10.46	0.75	25.51	14.27	5.83	23.80	3.72	43.10	16.90	...	1.82	8.47	24.61	4

5 rows × 3000 columns

```
In [12]: #some features are skewed and some are not, vary dependent on genes
import matplotlib

params = {'axes.titlesize':'12',
          'xtick.labelsize':'12',
          'ytick.labelsize':'12'}
matplotlib.rcParams.update(params)
data[features[0:10]].hist(bins=20,figsize=(30, 5),layout=(2,5))
```

```
Out[12]: array([[<AxesSubplot:title=('center': 'ENSG00000283050')>,<AxesSubplot:title=('center': 'ENSG00000178971')>,<AxesSubplot:title=('center': 'ENSG00000267530')>,<AxesSubplot:title=('center': 'ENSG00000181274')>,<AxesSubplot:title=('center': 'ENSG00000137414')>],<AxesSubplot:title=('center': 'ENSG00000268621')>,<AxesSubplot:title=('center': 'ENSG00000175309')>,<AxesSubplot:title=('center': 'ENSG00000103047')>,<AxesSubplot:title=('center': 'ENSG00000129226')>,<AxesSubplot:title=('center': 'ENSG00000133818')>]], dtype=object)
```



Q4: Summary of training at least three different classifier models, preferably of different nature in explainability and predictability. For example, you can start with a simple logistic regression as a baseline, adding other models or ensemble models. Preferably, all your models use the same training and test splits, or the same cross-validation method.

Answer:

Logistic regression, K-nearest neighbours, SVM, and decision tree are applied
Given only few data record, bagging and boosting may be less suit for this analysis

Preprocessing

```
In [13]: #Load preprocessing libraries
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelBinarizer


In [14]: #subset of data for standardize and binarize Labels
X=pd.DataFrame(StandardScaler().fit_transform(data[features]),columns=features)
Y=pd.DataFrame(LabelBinarizer().fit_transform(data['Tissue']),columns=['Tissue'])
print(X.iloc[:5,:4])
print(Y.head()) #C: 0 ; P:1
```

	ENSG00000283050	ENSG00000178971	ENSG00000267530	ENSG00000181274
0	1.081370	-0.381823	0.286446	2.284380
1	-1.679139	0.113386	1.687449	-0.659688
2	0.445119	-0.146172	0.448927	2.687876
3	1.040322	-1.187818	-0.332401	-0.793558
4	0.373284	0.615425	0.125619	0.616793

Tissue
0
1
2
3
4

```
In [90]: #data distribution after standard scaling, some skewness presist
X.iloc[:, :10].hist(bins=20, figsize=(30, 5), layout=(2, 5))
```

```
Out[90]: array([[<AxesSubplot:title=('center': 'ENSG00000283050')>,<AxesSubplot:title=('center': 'ENSG00000178971')>,<AxesSubplot:title=('center': 'ENSG00000267530')>,<AxesSubplot:title=('center': 'ENSG00000181274')>,<AxesSubplot:title=('center': 'ENSG00000137414')>],<AxesSubplot:title=('center': 'ENSG00000268621')>,<AxesSubplot:title=('center': 'ENSG00000175309')>,<AxesSubplot:title=('center': 'ENSG00000103047')>,<AxesSubplot:title=('center': 'ENSG00000129226')>,<AxesSubplot:title=('center': 'ENSG00000133818')>]], dtype=object)
```



```
In [27]: #Load libraries for classifier and reports
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegressionCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier

#Load report libraries
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score

#Grid search for best parameters
from sklearn.model_selection import GridSearchCV

import time
```

```
In [28]: #train-test split, double check for splitted data set
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.4, random_state=5)
print("X_train: {}, Y_train: {}, X_test: {}, Y_test: {}".format(X_train.shape, Y_train.shape, X_test.shape, Y_test.shape))

X_train: (24, 3000), Y_train: (24, 1), X_test: (16, 3000), Y_test: (16, 1),
```

Logistic Regression

```
In [81]: #Linear regression model
logR=GridSearchCV(estimator=LogisticRegression(solver='saga', n_jobs=-1),
                  param_grid={'cv':[3,5], 'penalty':['l1', 'l2', 'elasticnet']})

#train
logR_time = time.time()
logR.fit(X_train, Y_train)
logR_time_train = time.time() - logR_time

#test
logR_time = time.time()
Y_pred_logR = logR.predict(X_test)
logR_time_pred = time.time() - logR_time
```

KNN

```
In [82]: #KNN model
knn_time = time.time()
knn=GridSearchCV(estimator=KNeighborsClassifier(n_jobs=-1),
                  param_grid={'n_neighbors':[3,8]})

#train
knn_time = time.time()
knn.fit(X_train, Y_train)
knn_time_train = time.time() - knn_time

#test
knn_time = time.time()
Y_pred_knn = knn.predict(X_test)
knn_time_pred = time.time() - knn_time
```

SVM

```
In [83]: #SVM model
svm_time = time.time()
svc=GridSearchCV(estimator=SVC(),
                  param_grid={'C':[1e-3,1e-2,1e-1,1], 'kernel':['linear', 'poly', 'rbf'], 'degree':range(2,5)})

#train
svm_time = time.time()
svc.fit(X_train,Y_train)
svm_time_train = time.time() - svm_time
#test
knn_time = time.time()
Y_pred_svm = svc.predict(X_test)
svm_time_pred = time.time() - svm_time
```

Decision Tree

```
In [84]: #Decision tree model
dTree_time = time.time()
dTree=GridSearchCV(estimator=DecisionTreeClassifier(random_state=13),
                   param_grid={'criterion':['gini', 'entropy', 'log_loss'], 'max_depth':range(5,50,5)})

#train
dTree_time = time.time()
dTree.fit(X_train,Y_train)
dTree_time_train = time.time() - dTree_time
#test
dTree_time = time.time()
Y_pred_dTree = dTree.predict(X_test)
dTree_time_pred = time.time() - dTree_time
```

Analysis of each classifiers

```
In [85]: score_df = pd.DataFrame()

for i, j in enumerate([Y_pred_logR, Y_pred_knn, Y_pred_svm, Y_pred_dTree]):
    score_df[i] = [accuracy_score(Y_test,j),
                  precision_score(Y_test,j),
                  recall_score(Y_test,j),
                  f1_score(Y_test,j),
                  roc_auc_score(Y_test,j)]

score_df = score_df.T
score_df.columns=['accuracy','precision','recall','f1','auc']
score_df.index=['Logistic Regression','KNN','SVM','Decision Tree']
score_df['Training time']=[logR_time_train, knn_time_train, svm_time_train, dTree_time_train]
score_df['Prediction time']=[logR_time_pred, knn_time_pred, svm_time_pred, dTree_time_pred]
score_df
```

```
Out[85]:
```

	accuracy	precision	recall	f1	auc	Training time	Prediction time
Logistic Regression	1.0000	1.0	1.0	1.000000	1.00	21.612164	0.020988
KNN	1.0000	1.0	1.0	1.000000	1.00	0.835522	0.056967
SVM	0.9375	1.0	0.9	0.947368	0.95	8.081401	8.103387
Decision Tree	0.9375	1.0	0.9	0.947368	0.95	7.195887	0.020989

```
In [86]: # Display parameters for each estimators
print("Logistic Regression: ",logR.best_estimator_)
print("KNN: ",knn.best_estimator_)
print("SVM: ",svc.best_estimator_)
print("Decision Tree: ",dTree.best_estimator_)

Logistic Regression: LogisticRegressionCV(cv=3, n_jobs=-1, penalty='l1', solver='saga')
KNN: KNeighborsClassifier(n_jobs=-1, n_neighbors=3)
SVM: SVC(C=1, degree=2)
Decision Tree: DecisionTreeClassifier(max_depth=5, random_state=13)
```

```
In [87]: #visualization
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```


```
In [88]: #Y_test data
Test_Data=pd.DataFrame(Y_test.value_counts().to_list(),index=['Para','Cancer'],columns=['Counts'])
Test_Data
```

```
Out[88]:
```

	Counts
Para	10
Cancer	6

```
In [89]: #plot curves for each classification methods
from sklearn.metrics import confusion_matrix
sns.set_context('talk')

fig, axes = plt.subplots(1, 4, figsize=(26,5))
for i, j in enumerate([Y_pred_logR, Y_pred_knn, Y_pred_svm, Y_pred_dTree]):
    cm = confusion_matrix(Y_test, j)
    sns.heatmap(cm, ax=axes[i], annot=True, fmt='d', annot_kws={"size": 24, "weight": "bold"})
    axes[i].set_title(score_df.index[i])
    axes[i].set_xticklabels(['Cancer','Para'], fontsize=14);
    axes[i].set_yticklabels(['Cancer','Para'], fontsize=14);
    axes[i].set_ylabel('Ground Truth', fontsize=18)
    axes[i].set_xlabel('Prediction', fontsize=18);
```



Q5: A paragraph explaining which of your classifier models you recommend as a final model that best fits your needs in terms of accuracy and explainability.

Answer:

For this data set, all 4 tested classifiers showed strong predictability.

Among which, Logistic Regression and KNN perform best (highest accuracy, precision, and recall)

Comparatively, KNN with 3 nearest neighbours perform fast in both training and prediction, therefore may best suit for this data set, which contain large number of features (genes)

Q6: Summary Key Findings and Insights, which walks your reader through the main drivers of your model and insights from your data derived from your classifier model.

Answer:

Given large amount of features granted the stronger predictability for all tested models: Logistic Regression, KNN, SVM, and Decision Tree.

Several runs were performed, and number of random selected features/genes from 500 to 5000 were tested. The outcomes turned out to be almost identical. This implies that given high amount of features, individual feature may not play a decisive role in classification models. These also suggest the robustness of this model.

Q7: Suggestions for next steps in analyzing this data, which may include suggesting revisiting this model after adding specific data features that may help you achieve a better explanation or a better prediction.

Answer:

Given the KNN model is highly predictable for this data set, more data sets of similar processing steps may be applied for classification. Given the robustness(100% predictability) of this model, more gene expression profiles may be used to test the predictability of this model.