

Differentiate Colorectal carcinoma and para-carcinoma tissue based on expression profile

Title: Hypoxia-induced cysteine metabolism reprogramming are crucial for the tumorigenesis of colorectal cancer

Gene Expression Omnibus: GSE223119

Data Analyzer: Kai-Wei Chang

Data Owner: Shiyi Yang

Email: octoberfirst@sjtu.edu.cn

Affiliation: Shanghai Jiao Tong University, Shanghai, China

Q1: Main objective of the analysis that specifies whether your model will be focused on prediction or interpretation.

Answer:

This analysis will be focused on prediction

By determining whether the gene expression profiles capable of classifying the sample's origin, future expression profiles of unknown origin may be classified accordingly

In addition, the analysis will compare and contrast the performance of classification tools on low sample, high-dimension data set.

ie. each gene is assumed to be independent feature for this analysis

Q2: Brief description of the data set you chose and a summary of its attributes

Answer:

Metabolic reprogramming is a hallmark of human cancer and cancer-specific metabolism provide opportunities for cancer diagnosis, prognosis, and treatment. However, how metabolic pathways affect the initiation and progression of colorectal cancer remain largely unknown.

This data set includes 40 gene expression profiles of cancer and para-carcinoma tissue (tissue surrounding cancer).

Raw data is composed of gene ID, gene expression on each sample, gene names, and database and gene network category

This analysis attempt to build a model for predicting the tissue origin of unknown gene profiles

```
In [1]: ##No warning output
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn
```

Q3: Brief summary of data exploration and actions taken for data cleaning and feature engineering.

Answer:

Data Exploration:

- Gene expression has been preprocessed by the data provider, with no missing values
- This data set contain extremely high amount of features (ie. 61,700 genes)

Data Engineering:

- For the project purpose, only gene ID (feature), gene expression, and tissue origin will be considered:
- Removing features that are not expressing or have very low variability among all profiles
- For simplification in this analysis, I randomly selected 3000 features for analysis
- Gene expression may or may not be skewed, and expression levels vary from gene to gene, standardization is applied

Detail analysis see below

```
In [2]: #Import Data processing libraries
import numpy as np
import pandas as pd
```

```
In [3]: #Load Data, data is preprocessed
rawData=pd.read_csv(r"C:\Users\kai-w\Desktop\03_Supervised Machine Learning Classification\GSE223119_MJ20190424016-gene.tpm.matrix.annot.txt',sep='\t',header=0)
```

```
In [4]: #Data overview
rawData.head()
```

Out[4]:	Unnamed: 0	C10	C11	C12	C13	C14	C15	C18	C1	C20	...	cog	cog_description	KO id	KO_name	paths	pfam	go	n
0	ENSG000000000003	61.65	37.27	85.58	68.76	32.17	49.63	64.18	10.96	36.86	...	ENOG4111IRY(S-Function unknown)	ENOG4111IRY(Tenomodulin)	K17295	TSPAN6	NaN	PF00335.17(Tetraspannin: Tetraspanin family)	GO:0039532(biological_process:negative regulat...	NP_003261.1(tetraspanin-I isoform a [Homo sapi...
1	ENSG000000000005	0.31	3.39	1.34	2.24	0.25	4.48	1.64	0.15	0.89	...	ENOG410YB96(S-Function unknown)	ENOG410YB96(Tenomodulin)	NaN	NaN	NaN	PF04089.11(BRICHOS:BRICHOS domain)	GO:0005737(cellular_component:cytoplasm); GO:0...	XP_006986474.1(PREDICTED tenomodulin [Peromys...
2	ENSG0000000000419	81.96	34.78	56.60	81.13	49.78	121.31	91.79	39.23	41.02	...	COG0463(M-Cell wall/membrane/envelope biogenesis)	COG0463(Glycosyl transferase, family 2)	K00721	DPM1	map00510(N-Glycan biosynthesis)	PF00535.23(Glycos_transf_2-Glycosyl transferas...	GO:0019673(biological_process:GDP-mannose meta...	NP_001303964.1(dolichol phosphate mannosyltran...
3	ENSG0000000000457	4.98	1.42	2.84	4.01	4.90	3.86	3.63	4.08	3.95	...	ENOG410XQTGS(Function unknown)	ENOG410XQTGS(cerevisiae)	K17542	SCYL3	NaN	PF00069.22(Pkinase:Protein kinase domain); PF0...	GO:0005794(cellular_component:Golgi apparatus)...	XP_003893590.2(protein associating with the ca...
4	ENSG0000000000460	7.14	1.99	2.80	9.40	5.90	4.13	9.92	4.05	4.21	...	ENOG4110VTC(S-Function unknown)	ENOG4110VTC(Chromosome 1 open reading frame 112)	NaN	NaN	NaN	PF14868.3(DUF4487:Domain of unknown function (...)	NaN	XP_005245374.1(uncharacterize protein C1orf11...
5 rows x 54 columns																			

```
In [5]: #Data types for each column
rawData.dtypes
```

```
Out[5]: Unnamed: 0      object
C10      float64
C11      float64
C12      float64
C13      float64
C14      float64
C15      float64
C18      float64
C1       float64
C20      float64
C22      float64
C23      float64
C24      float64
C25      float64
C3       float64
C4       float64
C5       float64
C6       float64
C7       float64
C8       float64
C9       float64
P10      float64
P11      float64
P12      float64
P13      float64
P14      float64
P15      float64
P18      float64
P1       float64
P20      float64
P22      float64
P23      float64
P24      float64
P25      float64
P3       float64
P4       float64
P5       float64
P6       float64
P7       float64
P8       float64
P9       float64
gene_name object
length    int64
description object
cog        object
cog_description object
K0_id      object
K0_name    object
paths      object
pfam       object
go         object
nr         object
swissprot  object
entrez     float64
dtype: object

In [6]: #Transpose so treat genes as feature, and each profile as independent record
data=rawData.loc[:,rawData.columns.str.match('^C|^P')].rename(index=rawData['Unnamed: 0']).T
data.head()
```

	ENSG000000000003	ENSG000000000005	ENSG000000000419	ENSG000000000457	ENSG000000000460	ENSG000000000938	ENSG000000000971	ENSG000000001036	ENSG000000001084	ENSG000000001167	...	MSTRG.9662	MSTRG.9664	MSTRG.9666	MSTRG.9667	MSTRG.9925
C10	61.65	0.31	81.96	4.98	7.14	4.40	15.56	38.45	18.87	12.65	...	0.02	0.00	0.0	0.13	0.37
C11	37.27	3.39	34.78	1.42	1.99	3.10	5.15	34.72	11.98	3.12	...	0.10	6.57	0.0	1.17	0.61
C12	85.58	1.34	56.60	2.84	2.80	5.31	8.48	29.41	25.83	11.12	...	0.07	0.00	0.0	0.20	0.17
C13	68.76	2.24	81.13	4.01	9.40	4.46	14.99	45.74	25.93	13.76	...	0.04	0.00	0.0	0.33	0.75
C14	32.17	0.25	49.78	4.90	5.90	2.46	10.40	33.28	17.18	11.91	...	0.27	7.07	0.0	0.36	0.54

5 rows × 61700 columns

```
In [7]: #check for presence of empty values
data.isna().values.sum()

Out[7]: 0

In [8]: #huge feature size, reduce by taking randomly 1000 for the purpose of the project
data.shape

Out[8]: (40, 61700)

In [9]: #some features/genes have low variability, remove those
feature_var=pd.DataFrame([[[i, data[i].std()] for i in data.columns[1:11]],
                           columns=['feature', 'std']].set_index('feature'))
print(feature_var[feature_var['std']>0].sort_values('std'))

          std
feature
ENSG000000228856    0.001581
ENSG000000231051    0.001581
ENSG000000156925    0.001581
ENSG000000250231    0.001581
ENSG000000232264    0.001581
...
ENSG000000211890    6428.959034
ENSG000000240040    6780.842254
ENSG000000212907    7435.193295
ENSG000000198804    8216.179258
ENSG000000228253    18410.117553

[41927 rows x 1 columns]

In [10]: #random select 3000 features/genes that are variable for the project's purpose
features=feature_var[feature_var['std']>0.5].index.to_series().sample(3000).to_list()
data[features].head()
```


	ENSG00000116731	ENSG00000105379	ENSG00000026908	ENSG00000076356	ENSG000000234147	ENSG00000137770	ENSG000000280721	ENSG000000100605	ENSG000000168065	ENSG00000177383	...	ENSG000000206069	ENSG00000102996	ENSG00000177731	ENSG000000
C10	8.25	95.22	3.32	9.67	0.00	8.97	1.90	32.16	0.40	26.98	...	1.32	25.51	95.05	
C11	5.14	151.07	0.31	4.69	2.78	1.46	3.64	48.93	0.19	19.29	...	4.92	31.44	124.75	
C12	8.53	67.01	0.29	8.34	0.46	5.27	1.54	44.18	3.34	17.14	...	0.85	34.43	75.61	
C13	12.99	48.89	1.45	5.64	0.40	7.68	1.76	42.32	1.24	19.71	...	3.04	30.77	82.29	
C14	14.51	108.77	1.52	8.08	0.00	7.86	0.90	50.16	0.32	27.59	...	3.10	33.86	104.32	

5 rows × 3000 columns

```
In [11]: #some features are skewed and some are not, vary dependent on genes
import matplotlib

params = {'axes.titlesize':'12',
          'xtick.labelsize':'12',
          'ytick.labelsize':'12'}
matplotlib.rcParams.update(params)
data[features[0:10]].hist(bins=20,figsize=(30, 5),layout=(2,5))

Out[11]: array([[<AxesSubplot:title=('center': 'ENSG00000116731')>,
<AxesSubplot:title=('center': 'ENSG00000105379')>,
<AxesSubplot:title=('center': 'ENSG00000026908')>,
<AxesSubplot:title=('center': 'ENSG00000076356')>,
<AxesSubplot:title=('center': 'ENSG000000234147')>],
[<AxesSubplot:title=('center': 'ENSG00000137770')>,
<AxesSubplot:title=('center': 'ENSG000000280721')>,
<AxesSubplot:title=('center': 'ENSG000000100605')>,
<AxesSubplot:title=('center': 'ENSG000000168065')>,
<AxesSubplot:title=('center': 'ENSG00000177383')>]], dtype=object)
```



Q4: Summary of training at least three linear regression models which should be variations that cover using a simple linear regression as a baseline, adding polynomial effects, and using a regularization regression. Preferably, all use the same training and test splits, or the same cross-validation method.

Answer:

Linear Regression, Polynomial Regression, Linear Regression with LASSO regularization
Given only few data record, bagging and boosting may be less suit for this analysis

Based on the findings by the research group, they found the transporter genes of cystine and cysteine are all upregulated in colorectal cancer by tumor microenvironment induced ROS through transcription factor ATF4. Given such fact, whether gene expression pattern may be used to predict ATF4 expression pattern is of interest.

Preprocessing

```
In [12]: #Load preprocessing libraries
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import PolynomialFeatures

In [13]: #Identify ID that map to ATF4
ATF4_id=[rawData[rawData_gene_name=="ATF4"].iloc[0,0]]

In [14]: #subset of data for standardize and binarize labels
X=pd.DataFrame(StandardScaler().fit_transform(data[features]),columns=features)
Y=pd.DataFrame(StandardScaler().fit_transform(data[ATF4_id]),columns=[ 'ATF4' ])
pf = PolynomialFeatures(degree=2, include_bias=False)
X_pf = pf.fit_transform(data[features])
print(X.iloc[:5,:4])
print(X_pf)
print(Y.head())

      ENSG00000116731  ENSG00000105379  ENSG00000260908  ENSG00000076356
0      -0.447179      -0.006876      0.855265      0.264158
1      -1.609993      2.144600      -0.977947      -1.429403
2      -0.342488      -1.093593      -0.984812      -0.188144
3      1.325085      -1.791619      -0.586631      -1.106336
4      1.893406      0.515102      -0.562603      -0.276563
[[8.2500000e+00 9.522000e+01 3.320000e+00 ... 5.2787600e+03
 0.000000e+00 0.000000e+00]
[5.140000e+00 1.510790e+02 3.100000e-01 ... 2.329992e+03
 0.000000e+00 0.000000e+00]
[8.530000e+00 6.701000e+01 2.900000e-01 ... 3.034908e+03
 0.000000e+00 0.000000e+00]
...
[8.950000e+00 9.145000e+01 2.910000e+00 ... 4.769283e+03
 4.627020e+01 4.489000e-01]
[1.034000e+01 8.020000e+01 9.350000e+00 ... 5.065168e+03
 0.000000e+00 0.000000e+00]
[3.540000e+00 1.141900e+02 1.510000e+00 ... 1.098980e+03
 9.909000e+00 9.000000e-02]]
ATF4
0 -0.110166
1  0.197281
2 -0.185572
3  0.536964
4  1.441271

In [15]: #data distribution after standard scaling, some skewness persist
X.iloc[:,5].hist(bins=20,figsize=(30, 3),layout=(1,5))

Out[15]: array([[<AxesSubplot:title='center': 'ENSG00000116731'>,
<AxesSubplot:title='center': 'ENSG00000105379'>,
<AxesSubplot:title='center': 'ENSG00000260908'>,
<AxesSubplot:title='center': 'ENSG00000076356'>,
<AxesSubplot:title='center': 'ENSG00000234147'>]], dtype=object)
```

```
In [16]: #train test split
from sklearn.model_selection import train_test_split

#Load libraries for classifier and reports
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso

#Load report libraries
from sklearn.metrics import r2_score

import time

In [17]: #train-test split, double check for splitted data set
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.4, random_state=5)
print("X_train: {}, Y_train: {}, X_test: {}, Y_test: {}".format(X_train.shape,Y_train.shape,X_test.shape,Y_test.shape) )

X_train: (24, 3000), Y_train: (24, 1), X_test: (16, 3000), Y_test: (16, 1),

Linear Regression

In [23]: #Linear regression model
linR=LinearRegression()

#train
linR_time = time.time()
linR.fit(X_train,Y_train)
linR_time_train = time.time() - linR_time
#test
linR_time = time.time()
Y_pred_linR = linR.predict(X_test)
linR_time_pred = time.time() - linR_time
linR.coef_

Out[23]: array([[ -4.73740817e-04,  5.03606954e-04, -3.72638149e-04, ...,
-4.22145505e-26,  8.06827903e-06,  1.20522737e-03]])

In [29]: #coefficients imply each feature takes little effects to the predictability of ATF4 expression
print("Number of coefficients greater than 0.01: ",(abs(linR.coef_) > 0.01).sum() )
print("Number of coefficients greater than 0.001: ",(abs(linR.coef_) > 0.001).sum() )

Number of coefficients greater than 0.01: 0
Number of coefficients greater than 0.001: 1221

LASSO

In [19]: #LASSO regularization
lso = Lasso(alpha=0.01) #try some levels of regularization

#train
lso_time = time.time()
lso.fit(X_train,Y_train)
lso_time_train = time.time() - lso_time
#test
lso_time = time.time()
Y_pred_lso = lso.predict(X_test)
lso_time_pred = time.time() - lso_time
```

```
#Identify features of importance
key_features=(lso.coef_ > 0)
print("Number of key features/genes may associate with ATF4 expression: ", key_features.sum())

Number of key features/genes may associate with ATF4 expression:  14

In [20]: #List genes may be of interact with ATF4
geneList=[]
for i in lso.feature_names_in_[key_features]:
    geneList = geneList + [rowData.gene_name[rowData.lloc[:,0].to_list().index(i)]]
geneList

Out[20]: ['AC011472.3',
'CATSPER2',
'TRIM26',
'GLRX2',
'YBX3',
nan,
'GTF2E2',
'TUBA8',
'UFD1',
'BST2',
'ZNF394',
'NME4',
'MCS1',
'AL365203.2']
```

Linear Regression with Polynomial Features

```
#train-test split with polynomial features
X_pf_train, X_pf_test, Y_pf_train, Y_pf_test = train_test_split(X_pf, Y, test_size=0.4, random_state=5)
print("X_train: {}, Y_train: {}, X_test: {}, Y_test: {}".format(X_pf_train.shape,Y_pf_train.shape,X_pf_test.shape,Y_pf_test.shape) )

#train
linR_pf = linearRegression()
linR_pf.time = time.time()
linR_pf.fit(X_pf_train,Y_pf_train)
linR_pf.time_train = time.time() - linR_pf.time
#test
linR_pf.time = time.time()
Y_pred_linR_pf = linR_pf.predict(X_pf_test)
linR_pf.time_pred = time.time() - linR_pf.time

X_train: (24, 4504500), Y_train: (24, 1), X_test: (16, 4504500), Y_test: (16, 1),
```

Analysis of each classifiers

```
In [22]: score_df = pd.DataFrame()

for i, j in enumerate([Y_pred_linR, Y_pred_lso, Y_pred_linR_pf]):
    score_df[i] = [r2_score(Y_test,j)]
score_df = score_df.T
score_df.columns=['R2']
score_df.index=['Linear Regression','Lasso','Polynomial']
score_df['Training time']=[linR_time_train, lso_time_train, linR_pf_time_train]
score_df['Prediction time']=[linR_time_pred, lso_time_pred, linR_pf_time_pred]
score_df
```

Out[22]:

	R2	Training time	Prediction time
Linear Regression	0.725526	0.033981	0.020987
Lasso	0.749997	0.154910	0.020989
Polynomial	0.383518	13.971140	0.216882

Q5: A paragraph explaining which of your regressions you recommend as a final model that best fits your needs in terms of accuracy and explainability.

Answer:
For this data set, Linear Regression and Lasso show similar predictability to ATF4 expression. Nevertheless, R-square of less than 0.75 suggest the predictability may not be very accurate in these tested models. On the other hand, polynomial features have dramatic reduced accuracy and high training and prediction time, therefore not recommended.

Q6: Summary Key Findings and Insights, which walks your reader through the main drivers of your model and insights from your data derived from your linear regression model.

Answer:
Comparing Linear regression and Lasso outcomes. Linear regression suggested each genes may take minor effects to ATF4 expression, therefore suggested a robust system in gene network, and that contradict with data provider's finding. On the other hand, Lasso regularization narrow down the effective terms to several genes, which may imply regulatory network for further analysis. Also these genes makes more sense in a gene regulatory model.

Q7: Suggestions for next steps in analyzing this data, which may include suggesting revisiting this model adding specific data features to achieve a better explanation or a better prediction.

Answer:
Several genes are involved in regulatory networks involvein ATF4, whether such predictability applicable for other genes of importance is of question. Additionally, the predictability is not effective, further analysis over the gene expresison may be required for better predictability based on gene expression profiles.

```
In [ ]:
```